

```

from transformers import pipeline
summarizer = pipeline("summarization")
text="""Hugging face is a company that specializes in natural language processing(NLP).it has developed the Transformers library,which pr
summary=summarizer(text,max_length=50,min_length=10,do_sample=False)
print(summary)

```

⚠ No model was supplied, defaulted to sshleifer/distilbart-cnn-12-6 and revision a4f8f3e (<https://huggingface.co/sshleifer/distilbart-cnn-12-6>)  
Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% 1.80k/1.80k [00:00<00:00, 78.5kB/s]

pytorch\_model.bin: 100% 1.22G/1.22G [00:17<00:00, 70.5MB/s]

model.safetensors: 89% 1.09G/1.22G [00:14<00:01, 113MB/s]

tokenizer\_config.json: 100% 26.0/26.0 [00:00<00:00, 627B/s]

vocab.json: 100% 899k/899k [00:00<00:00, 4.27MB/s]

merges.txt: 100% 456k/456k [00:00<00:00, 7.57MB/s]

Device set to use cpu

[{'summary\_text': ' Hugging face is a company that specializes in natural language processing . It has developed the Transformers l:

```

import torch
from transformers import AutoModelForSequenceClassification, AutoTokenizer
model_name = "cardiffnlp/tweet-topic-21-multi"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForSequenceClassification.from_pretrained(model_name)
labels = [
    "arts_&_culture", "business_&_entrepreneurs", "celebrity_&_pop_culture", "diaries_&_daily_life",
    "family", "fashion_&_style", "film_tv_&_video", "fitness_&_health", "food_&_dining",
    "gaming", "learning_&_educational", "music", "news_&_social_concern", "other_hobbies",
    "relationships", "science_&_technology", "sports_&_esports", "travel_&_adventure",
    "youth_&_student_life"
]

```



tokenizer_config.json: 100%	1.30k/1.30k [00:00<00:00, 14.8kB/s]
vocab.json: 100%	798k/798k [00:00<00:00, 6.49MB/s]
merges.txt: 100%	456k/456k [00:00<00:00, 8.83MB/s]
tokenizer.json: 100%	1.36M/1.36M [00:00<00:00, 16.6MB/s]
special_tokens_map.json: 100%	239/239 [00:00<00:00, 3.57kB/s]
config.json: 100%	1.88k/1.88k [00:00<00:00, 36.5kB/s]
pytorch_model.bin: 100%	499M/499M [00:08<00:00, 19.3MB/s]

```
import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification, AutoTokenizer

model_name="cardiffnlp/tweet-topic-21-multi"
tokenizer=AutoTokenizer.from_pretrained(model_name)
model=AutoModelForSequenceClassification.from_pretrained(model_name)

labels=[
    "arts_&culture", "business&entrepreneurs", "celebrity&pop_culture", "diaries&daily_life",
    "family", "fashion_&style", "film_tv&video", "fitness&health", "food&dining",
    "gaming", "learning_&educational", "music", "news&social_concern", "other_hobbies", "relationships",
    "science_&technology", "sports&esports", "travel&adventure", "youth&student_life"
]

texts=[
    "The latest iPhone was just released with an incredible new camera!",
    "Manchester United won their match with a stunning goal in the last minute."
    "NASA just launched a new mission to explore the surface of Mars.",
    "The Oscars had some surprising winners this year!"
]

inputs=tokenizer(texts, padding=True, truncation=True, return_tensors="pt")
with torch.no_grad():
    outputs=model(**inputs)
```

```
probabilities=torch.nn.functional.softmax(outputs.logits,dim=-1)
predictions=torch.argmax(probabilities,dim=1)
```

```
for text,pred,prob in zip(texts,predictions,probabilities):
    print(f"Text: {text}\nTopic: {labels[pred.item()]}, Confidence: {prob[pred].item():.4f}\n")
```

➞ Asking to truncate to max\_length but no maximum length is provided and the model has no predefined maximum length. Default to no truncate.  
 Text: The latest iPhone has just released with an incredible new camera!  
 Topic: science&technology, Confidence: 0.9260

Text: Manchester United won their match with a stunning goal in the last minute.NASA just launched a new mission to explore the surface of Mars.  
 Topic: sports&sports, Confidence: 0.7513


Text: The Oscars had some surprising winners this year!  
 Topic: film\_tv&video, Confidence: 0.9357

```
from transformers import AutoModelForCausalLM, AutoTokenizer
model_name = "gpt2"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(model_name)
prompt = "I love coding languages,"
inputs = tokenizer(prompt, return_tensors="pt")
output = model.generate(**inputs, max_length=50, num_return_sequences=1, temperature=0.7, top_k=50)
generated_text = tokenizer.decode(output[0], skip_special_tokens=True)
print(generated_text)
```

➞ Setting `pad\_token\_id` to `eos\_token\_id`:50256 for open-end generation.  
 I love coding languages, but I'm not a programmer. I'm a programmer. I'm a programmer. I'm a programmer. I'm a programmer. I'm a programmer.

```
prompt="Hello, I'm a language model"
inputs=tokenizer(prompt,return_tensors="pt")
output=model.generate(**inputs,max_length=50,num_return_sequences=1,temperature=0.7,top_k=50)
```

```
generated_text=tokenizer.decode(output[0],skip_special_tokens=True)
print(generated_text)
```

 /usr/local/lib/python3.11/dist-packages/transformers/generation/configuration\_utils.py:628: UserWarning: `do\_sample` is set to `False` but will be ignored because `num\_beams` is greater than 1. In this case, only the most likely next token is returned regardless of sampling parameters. To actually sample the next token, use a sampling strategy like `do\_sample=True, top\_k=1` or `do\_sample=True, top\_p=0.5`.  
warnings.warn(  
Setting `pad\_token\_id` to `eos\_token\_id`:50256 for open-end generation.  
Hello, I'm a language modeler. I'm a programmer. I'm a programmer. I'm a programmer. I'm a programmer. I'm a programmer. I'm a programmer.

```
import os
import atexit
import shutil
from transformers import BlenderbotTokenizer, BlenderbotForConditionalGeneration
model_name = "facebook/blenderbot-1B-distill"
tokenizer = BlenderbotTokenizer.from_pretrained(model_name)
model = BlenderbotForConditionalGeneration.from_pretrained(model_name)
def interact_with_chatbot(user_input, conversation_history):
    conversation_history.append(f"User: {user_input}")
    inputs = tokenizer(conversation_history, return_tensors="pt", padding=True, truncation=True)
    outputs = model.generate(**inputs, max_length=100, num_return_sequences=1)

    outputs = model.generate(**inputs, max_length=100, num_return_sequences=1)
```

\*\*\* tokenizer\_config.json: 100% 1.05k/1.05k [00:00<00:00, 23.8kB/s]

Start coding or [generate](#) with AI.

merges.txt: 100% 62.9k/62.9k [00:00<00:00, 1.47MB/s]

added\_tokens.json: 100% 16.0/16.0 [00:00<00:00, 399B/s]

special\_tokens\_map.json: 100% 772/772 [00:00<00:00, 25.0kB/s]

tokenizer.json: 100% 310k/310k [00:00<00:00, 4.18MB/s]

config.json: 100% 1.38k/1.38k [00:00<00:00, 37.0kB/s]

pytorch\_model.bin: 100% 2.87G/2.87G [00:53<00:00, 60.4MB/s]

model.safetensors: 40% 1.15G/2.87G [00:20<00:30, 56.9MB/s]