

```
!pip install pytextrank
```



```
Requirement already satisfied: spacy-legacy<3.1.0,>=3.0.11 in /usr/local/lib/python3
Requirement already satisfied: spacy-loggers<2.0.0,>=1.0.0 in /usr/local/lib/python3
Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in /usr/local/lib/python3.1
Requirement already satisfied: cymem<2.1.0,>=2.0.2 in /usr/local/lib/python3.11/dist
```

```
Requirement already satisfied: pillow<9.0.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-pa
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dis
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packag
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-pack
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.11/c
Requirement already satisfied: pydantic-core==2.27.2 in /usr/local/lib/python3.11/di
Requirement already satisfied: typing-extensions>=4.12.2 in /usr/local/lib/python3.1
```


Downloading colorama-0.4.6-py2.py3-none-any.whl (25 kB)

Downloading executing-2.2.0-py2.py3-none-any.whl (26 kB)

Installing collected packages: executing, colorama, asttokens, icecream, pytextrank

Successfully installed asttokens-3.0.0 colorama-0.4.6 executing-2.2.0 icecream-2.1.4


```
import spacy
import pytextrank
```

 /usr/local/lib/python3.11/dist-packages

```
document="""Not only did it only confirm that the film would be
unfunny and generic,but it also managed to give away the
ENTIRE movie; and I'm not exaggerating - every moment,every plot point,
every joke is told in the trailer."""
```


```
import en_core_web_sm
nlp = en_core_web_sm.load()
```

```
en_nlp = spacy.load("en_core_web_sm")
en_nlp.add_pipe("textrank")
doc = en_nlp(document)
```


 /usr/local/lib/python3.11/dist-packages/spacy/util.py:1740: UserWarning: [W111] Jupyter warnings.warn(Warnings.W111)



```
tr=doc._.textrank  
print(tr.elapsed_time)
```

 21.96216583251953

```
for combination in doc._.phrases:  
    print(combination.text,combination.rank,combination.count)
```

 ENTIRE 0.13514348101679782 1
the
 ENTIRE movie 0.06930139138062855 1
every joke 0.05936552514177136 1
the film 0.05423292745389326 1
the trailer 0.04834919915077192 1
I 0.0 1
it 0.0 2

```
from bs4 import BeautifulSoup  
from urllib.request import urlopen
```

```
def get_only_text(url):
    page=urlopen(url)
    soup=BeautifulSoup(page)
    text='\t'.join(map(lambda p:p.text,soup.find_all('p')))
    print(text)
    return soup.title.text,text
```

```
url="https://en.wikipedia.org/wiki/Natural_language_processing"
text=get_only_text(url)
```



Natural language processing (NLP) is a subfield of computer science and especially artificial intelligence. Major tasks in natural language processing are speech recognition, text classification, and machine translation. Natural language processing has its roots in the 1950s.[1] Already in 1950, Allen Newell and Herbert A. Simon developed the first NLP system, the General Problem Solver. The premise of symbolic NLP is well-summarized by John Searle's Chinese room experiment. Up until the 1980s, most natural language processing systems were based on computational linguistics, a symbolic approach, i.e., the hand-coding of a set of rules for manipulating symbols. Machine learning approaches, which include both statistical and neural networks, became dominant in the 1990s. Although rule-based systems for manipulating symbols were still in use in 2020, they were commonly used before that. In the late 1980s and mid-1990s, the statistical approach ended a period of AI winters. The earliest decision trees, producing systems of hard if-then rules, were still used. Only the introduction of hidden Markov models, applied to part-of-speech tagging, annotation, and machine translation, marked a major breakthrough. A major drawback of statistical methods is that they require elaborate feature engineering. Intermediate tasks (e.g., part-of-speech tagging and dependency parsing) are not solved directly. Neural machine translation, based on then-newly invented sequence-to-sequence models, became the dominant approach in the 2010s.

The following is a list of some of the most commonly researched tasks in natural language processing. Though natural language processing tasks are closely intertwined, they can be categorized into different groups. Based on long-standing trends in the field, it is possible to extrapolate future research directions. Most higher-level NLP applications involve aspects that emulate intelligent behavior. Cognition refers to "the mental action or process of acquiring knowledge and understanding." As an example, George Lakoff offers a methodology to build natural language processing systems. Ties with cognitive linguistics are part of the historical heritage of NLP, but

```
len(''.join(text))
```

→ 6587

```
text[:1000]
```

→ ('Natural language processing - Wikipedia',
'Natural language processing (NLP) is a subfield of computer science and especially artificial intelligence. It is primarily concerned with providing computers with the ability to process data encoded in natural language and is thus closely related to information retrieval, knowledge representation and computational linguistics, a subfield of linguistics. Typically data is collected in text corpora, using either rule-based, statistical or neural-based approaches in machine learning and deep learning.\n\tMajor tasks in natural language processing are speech recognition, text classification, natural-language understanding, and natural-language generation.\n\tNatural language processing has its roots in the

1950s.[1] Already in 1950, Alan Turing published an article titled "Computing Machinery and Intelligence" which proposed what is now called the Turing test as a criterion of intelligence, though at the time that was not articulated as a problem separate from artificial intelligence. The proposed test includes a task that involves the automated interpretation and generation of natural language.\n\tThe premise of symbolic NLP is well-summarized by John Searle\'s Chinese room experiment: Given a collection of rules (e.g., a Chinese phrasebook, with questions and matching answers), the computer emulates natural language understanding (or other NLP tasks) by applying those rules to the data it confronts.\n\tUp until the 1980s, most natural language processing systems were based on complex sets of hand-written rules. Starting in the late 1980s, however, there was a revolution in natural language processing with the introduction of machine learning algorithms for language processing. This was due to both the steady increase in computational power (see Moore\'s law) and the gradual lessening of the dominance of Chomskyan theories of linguistics (e.g. transformational grammar), whose theoretical underpinnings discouraged the sort of corpus linguistics that underlies the machine-learning approach to language processing.[8]\n\tSymbolic approach, i.e., the hand-coding of a set of rules for manipulating symbols, coupled with a dictionary lookup, was historically the first approach used both by AI in general and by NLP in particular:[18][19] such as by writing grammars or devising heuristic rules for stemming.\n\tMachine learning approaches, which include both statistical and neural networks, on the other hand, have many advantages over the symbolic approach: \n\tAlthough rule-based systems for manipulating symbols were still in use in 2020, they have become mostly obsolete with the advance of LLMs in 2023. \n\tBefore that they were commonly used:\n\tIn the late 1980s and mid-1990s, the statistical approach ended a period of AI winter, which was caused by the inefficiencies of the rule-based approaches.[20][21]\n\tThe earliest decision trees, producing systems of hard if-then rules, were still very similar to the old

rule-based approaches. Only the introduction of hidden Markov models, applied to part-of-speech tagging, announced the end of the old rule-based approach. A major drawback of statistical methods is that they require elaborate feature engineering. Since 2015,[22] the statistical approach has been replaced by the neural networks approach, using semantic networks[23] and word embeddings to capture semantic properties of words. Intermediate tasks (e.g., part-of-speech tagging and dependency parsing) are not needed anymore. Neural machine translation, based on then-newly invented sequence-to-sequence transformations, made obsolete the intermediate steps, such as word alignment, previously necessary for statistical machine translation. The following is a list of some of the most commonly researched tasks in natural language processing. Some of these tasks have direct real-world applications, while others more commonly serve as subtasks that are used to aid in solving larger tasks. Though natural language processing tasks are closely intertwined, they can be subdivided into categories for convenience. A coarse division is given below. Based on long-standing trends in the field, it is possible to extrapolate future directions of NLP. As of 2020, three trends among the topics of the long-standing series of CoNLL Shared Tasks can be observed:[46] Most higher-level NLP applications involve aspects that emulate intelligent behaviour and apparent comprehension of natural language. More broadly speaking, the technical operationalization of increasingly advanced aspects of

```
!pip install sumy
```



Collecting sumy

Downloading sumy-0.11.0-py2.py3-none-any.whl.metadata (7.5 kB)

Collecting docopt<0.7,>=0.6.1 (from sumy)

Downloading docopt-0.6.2.tar.gz (25 kB)

Preparing metadata (setup.py) ... done


```

Collecting breadability>=0.1.20 (from sumy)
  Downloading breadability-0.1.20.tar.gz (32 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: requests>=2.7.0 in /usr/local/lib/python3.11/dist-packages
Collecting pycountry>=18.2.23 (from sumy)
  Downloading pycountry-24.6.1-py3-none-any.whl.metadata (12 kB)
Requirement already satisfied: nltk>=3.0.2 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: chardet in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: lxml>=2.0 in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: click in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: joblib in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from nltk)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from nltk)
Downloading sumy-0.11.0-py2.py3-none-any.whl (97 kB)
----- 97.3/97.3 kB 6.0 MB/s eta 0:00:00
Downloading pycountry-24.6.1-py3-none-any.whl (6.3 MB)
----- 6.3/6.3 MB 21.0 MB/s eta 0:00:00
Building wheels for collected packages: breadability, docopt
  Building wheel for breadability (setup.py) ... done
  Created wheel for breadability: filename=breadability-0.1.20-py2.py3-none-any.whl size=13706 sha256=1a0b08c4b75c4116c31f83c8f9f047231251e9
  Stored in directory: /root/.cache/pip/wheels/4d/57/58/7e3d7fedf51fe248b7fcee3df6945a6
  Building wheel for docopt (setup.py) ... done
  Created wheel for docopt: filename=docopt-0.6.2-py2.py3-none-any.whl size=13706 sha256=1a0b08c4b75c4116c31f83c8f9f047231251e9
  Stored in directory: /root/.cache/pip/wheels/1a/b0/8c/4b75c4116c31f83c8f9f047231251e9

```

```
Successfully built breadability docopt
Installing collected packages: docopt, pycountry, breadability, sumy
Successfully installed breadability-0.1.20 docopt-0.6.2 pycountry-24.6.1 sumy-0.11.0
```

```
!pip install lxml-html-clean
from sumy.parsers.html import HtmlParser
from sumy.parsers.plaintext import PlaintextParser
from sumy.nlp.tokenizers import Tokenizer
from sumy.summarizers.lsa import LsaSummarizer as Summarizer
from sumy.nlp.stemmers import Stemmer
from sumy.utils import get_stop_words
from sumy.summarizers.lex_rank import LexRankSummarizer as Summarizer
from sumy.summarizers.lsa import LsaSummarizer
```



```
Requirement already satisfied: lxml-html-clean in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: lxml in /usr/local/lib/python3.11/dist-packages (from lxml-html-clean)
```

```
LANGUAGE="english"
SENTENCES_COUNT=10
url="https://en.wikipedia.org/wiki/Natural_language_processing"
parser=HtmlParser.from_url(url,Tokenizer(LANGUAGE))
summarizer=LsaSummarizer()
```

```

Summarizer=LsaSummarizer(Stemmer(LANGUAGE))
summarizer.stop_words=get_stop_words(LANGUAGE)
for sentence in Summarizer(parser.document,sentences_count=SENTENCES_COUNT):
    print(sentence)

```

➡ Generally, this task is much more difficult than supervised learning, and typically pro
 [25] For a language like English, this is fairly trivial, since words are usually separ
 The applications of sentiment analysis are diverse, extending to tasks such as categori
 Coreference resolution Given a sentence or larger chunk of text, determine which words
 Natural-language understanding(NLU) Convert chunks of text into more formal representat
 [44][45] Based on long-standing trends in the field, it is possible to extrapolate futu
 PMID 33736486.^ Lee, Jennifer; Yang, Samuel; Holland-Hall, Cynthia; Sezgin, Emre; Gill,
 Retrieved 5 December 2021.^ Yi, Chucai; Tian, Yingli(2012), "Assistive Text Reading fro
 Advances in Neural Information Processing Systems.^ Kariampuzha, William; Alyea, Gioco
 In some cases, there are few historical records on long-gone civilizations to serve as

```

import nltk
nltk.download('punkt_tab')

```

➡ [nltk_data] Downloading package punkt_tab to /root/nltk_data...
 [nltk_data] Unzipping tokenizers/punkt_tab.zip.
 True

```
text="""A vaccine for the coronavirus will likely be ready by early 2021 but rolling it out
a leading vaccine scientist told Bloomberg.India, which is host to some of the front-runne
currently has no local infrastructure in place to go beyond immunizing babies and pregnant
said Gagandeep Kang, professor of microbiology at the Vellore-based Christian Medical Coll
The timing of the vaccine is a contentious subject around the world.
```

```
In the U.S., President Donald Trump has contradicted a top administration health expert by
In India, Prime Minister Narendra Modi's government had promised an indigenous vaccine as e
a claim the government and its apex medical research body has since walked back."""
```

```
from sumy.parsers.plaintext import PlaintextParser
from sumy.nlp.tokenizers import Tokenizer
from sumy.summarizers.lex_rank import LexRankSummarizer
```

```
parser=PlaintextParser.from_string(text,Tokenizer("english"))
```

Generated code may be subject to a license | FenrirUnbound/too-lazy-decrease-reading

```
from sumy.summarizers.lex_rank import LexRankSummarizer
from sumy.utils import get_stop_words
summarizer_lex=LexRankSummarizer()
```

```
summarizer_lex.stop_words=get_stop_words('english')
summary =summarizer_lex(parser.document,5)
"
```

```
lex_summary=""  
for sentence in summary:  
    lex_summary+=str(sentence)+" "  
print(lex_summary)
```

⇒ A vaccine for the coronavirus will likely be ready by early 2021 but rolling it out sat



Start coding or generate with AI.